

# Asylum Seekers Entering the EU

## 1. Exploratory Data Analysis

---

### 1.1. Data Source, Headers and Labels

---

**DATASET** “Asylum and first time asylum applicants by citizenship, age and sex Monthly data (rounded) [migr\_asyappctzm]”

**LAST UPDATE\*** 19.03.19 06:51:45

**EXTRACTION DATE** 19.03.19 13:19:35

**SOURCE OF DATA** Eurostat

---

Attribute	Values
TIME	“2008M01” to “2018M12”
GEO	“European Union - 28 countries”
CITIZEN	“Extra-EU28”
SEX	“Total”, “Males”, “Females”
AGE	“Total”
ASYL_APP	“Asylum applicant”, “First time applicant”
UNIT	“Person”

---

No footnotes available.

Available flags:

- b, “break in time series”
  - c, “confidential”
  - d, “definition differs, see metadata”
  - e, “estimated”
  - f, “forecast”
  - n, “not significant”
  - p, “provisional”
  - r, “revised”
  - s, “Eurostat estimate”
  - u, “low reliability”
  - z, “not applicable”
  - Special value: “:”, “not available”
- 

### 1.2 Reading the data

```
#set working directory
setwd("~/DIT/Time Series 2 Forecasting/Project/2-Trend and seasonality")

#Graphical Parameters: For colours, color specifications, check colors() or, even better, demo(colors)
```

```

library(readr)
library(forecast)
library(tseries)
library(mgcv)

## Loading required package: nlme

##
## Attaching package: 'nlme'

## The following object is masked from 'package:forecast':
##
##      getResponse

## This is mgcv 1.8-26. For overview type 'help("mgcv-package")'.

datafile <- read_csv("migr_asylum_applicant_EU28/migr_asyappctzm_1_Data.csv")

```

```

## Parsed with column specification:
## cols(
##   TIME = col_character(),
##   DATE_PATTERN = col_character(),
##   GEO = col_character(),
##   CITIZEN = col_character(),
##   SEX = col_character(),
##   AGE = col_character(),
##   ASYL_APP = col_character(),
##   UNIT = col_character(),
##   Value = col_number(),
##   `Flag and Footnotes` = col_character()
## )

```

```
head(datafile)
```

```

## # A tibble: 6 x 10
##   TIME DATE_PATTERN GEO CITIZEN SEX AGE ASYL_APP UNIT Value
##   <chr> <chr>      <chr> <chr> <chr> <chr> <chr> <chr> <dbl>
## 1 2008~ 01/01/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 18355
## 2 2008~ 01/02/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 16885
## 3 2008~ 01/03/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 15570
## 4 2008~ 01/04/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 16975
## 5 2008~ 01/05/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 17110
## 6 2008~ 01/06/08 Euro~ Extra~~ Total Total Asylum ~ Pers~ 17520
## # ... with 1 more variable: `Flag and Footnotes` <chr>

```

All column names:

```
names(datafile)
```

```

## [1] "TIME"           "DATE_PATTERN"   "GEO"
## [4] "CITIZEN"        "SEX"            "AGE"
## [7] "ASYL_APP"       "UNIT"           "Value"
## [10] "Flag and Footnotes"

```


Convert the values into a time series Plot Values against Time

```

values = datafile[9]
values = ts(values, start=c(2008, 1), end=c(2018, 12), frequency=12)
ts.plot(values, main="Asylum Applications to EU28 Countries 2008-2018", ylab="Number of Applications")

```

## Asylum Applications to EU28 Countries 2008-2018



```
summary(values)
```

```
##      Value  
## Min.   : 15570  
## 1st Qu.: 22881  
## Median : 35925  
## Mean   : 48345  
## 3rd Qu.: 60488  
## Max.   :171895
```

```
end(values)
```

```
## [1] 2018 12
```

## 2. Smoothing the data

We can use smoothing to reduce the volatility in our observed data and making it into a more stable and predictable series.

### Trend Analysis - non-parametric

```
#Create equally spaced time points for fitting trends  
time.pts = c(1:length(values))  
time.pts = c(time.pts - min(time.pts))/max(time.pts)
```

## Fitting a moving average, kernel smoothing, loess and splines smoothing

```
#define mav/smoothing methods, and fit
values.mafilter.fit = filter(values, filter = rep(1/4, 4), sides = 2)

ma.fit = ma(values, order=2, centre=TRUE)
values.fit.ma = ts(ma.fit, start=c(2008, 1), frequency=12)

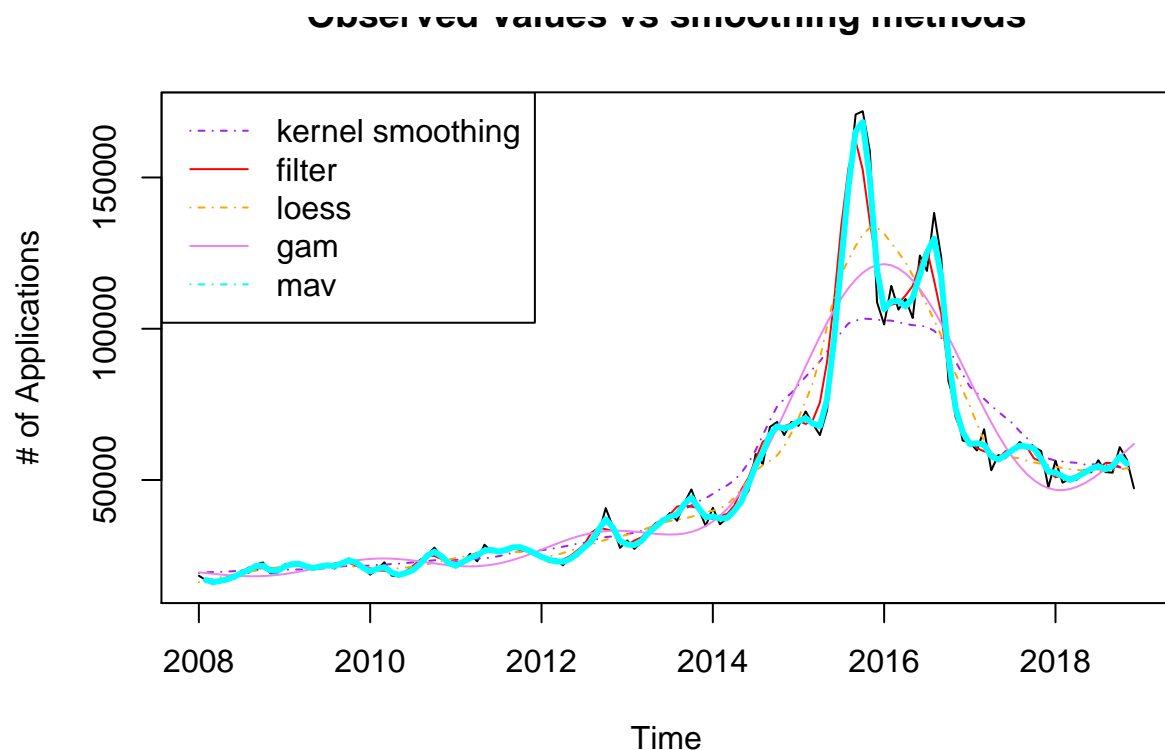
ksmooth.fit = ksmooth(time.pts, values, kernel = "box", bandwidth = 0.2)
values.fit.ksmooth = ts(ksmooth.fit$y, start=c(2008, 1), frequency=12)

loess.fit = loess(as.matrix(values)~time.pts, data=values, span=0.2)
values.fit.loess = ts(predict(loess.fit), start=c(2008, 1), frequency=12)

gam.fit = gam(values~s(time.pts))
values.fit.gam = ts(fitted(gam.fit), start=c(2008, 1), frequency=12)
```

## Plotting

```
# plot fits against values
ts.plot(values, ylab="# of Applications", main="Observed Values vs smoothing methods")
lines(values.fit.ksmooth, lwd=1, lty=4, col="purple")
lines(values.mafilter.fit, col="red")
lines(values.fit.loess, col="orange", lty=4)
lines(values.fit.gam, col="violet")
lines(values.fit.ma, col="cyan", lwd=3)
legend(x="topleft", c("kernel smoothing", "filter", "loess", "gam", "mav"), lty = c(4, 1), col=c("purple", "red", "orange", "violet", "cyan"))
```



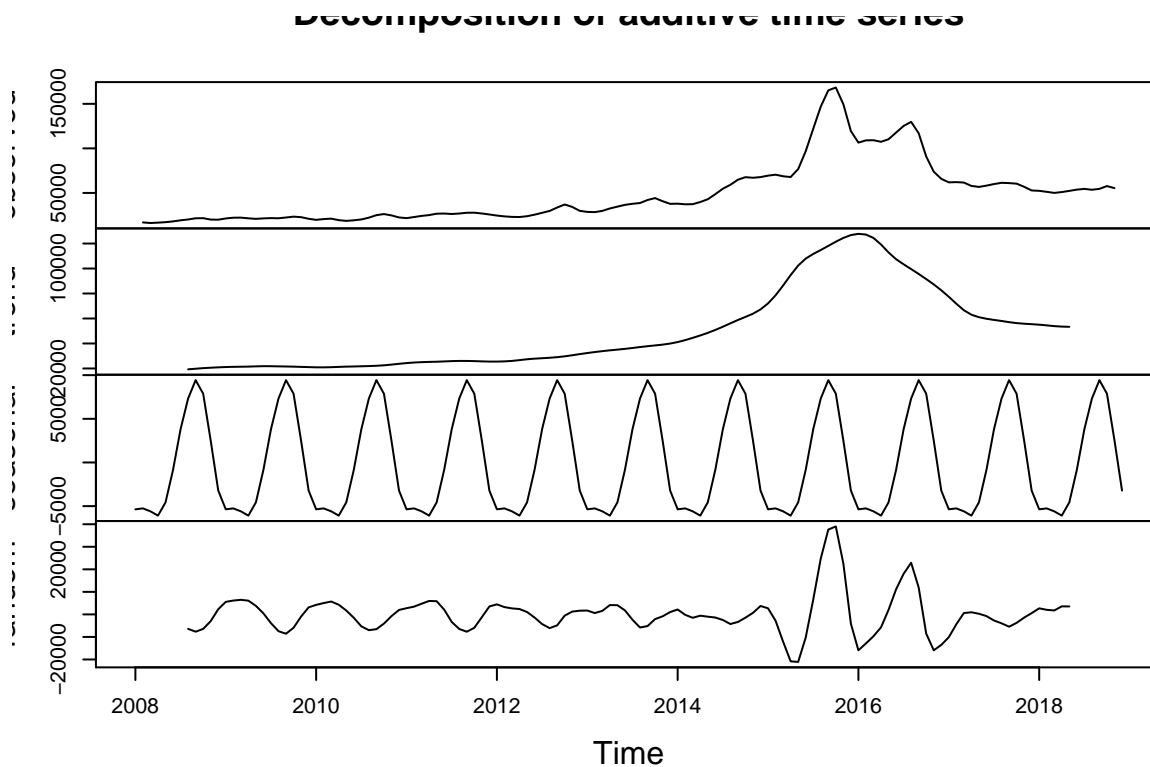
```
#values.fit.mav is the ma dataframe (type of the objects is float) with the transformed values for each
#ablines is an a, b line graphing function, a is y intercept and b is slope
```

It's a matter of balancing smoothness to make predictability easier and accuracy. Pick `filter(red)` or `moving average(cyan)`.

### 3. Decomposing the data


Trend component, seasonal component and residuals (random/white noise). We assume this series is additive (instead of multiplicative).

```
values.ma.decomp = decompose(values.fit.ma, type=c("additive"))  
plot(values.ma.decomp)
```




A more detailed decomposition of this time-series can be seen below by using STL - Seasonal Decomposition of Time Series by LOESS:

```
values.decomp.stl <- stl(na.omit(values.fit.ma[, 1]), s.window="periodic")  
plot(values.decomp.stl)
```



Using the decomposed time series, we can subtract the \$seasonal component from the data and de-seasonalize it - which can be useful to apply an ARIMA process, instead of SARIMA.

```
values.deseason = seasadj(values.decomp.stl)
plot(values.deseason, main="Deseasonalized time series")
```



```
print(values.decomp.stl)
```

```
## Call:
## stl(x = na.omit(values.fit.ma[, 1]), s.window = "periodic")
##
## Components
##      seasonal      trend      remainder
## Feb 2008 -5025.1317  19144.14   2804.74388
## Mar 2008 -5391.2925  19102.97   2538.32444
## Apr 2008 -5792.7941  19061.80   3388.49574
## May 2008 -4342.8346  19158.09   2363.49072
## Jun 2008 -747.0787   19254.39  -446.06072
## Jul 2008 3294.3252   19455.33 -3479.65891
## Aug 2008 6789.0248   19656.28 -6225.30284
## Sep 2008 8908.2448   19862.23 -7219.22019
## Oct 2008 7890.9851   20068.17 -6281.65786
## Nov 2008 2875.5202   20501.07 -3276.59209
## Dec 2008 -3149.1758  20933.97   2185.20483
## Jan 2009 -5309.7893  21417.41   5231.12483
## Feb 2009 -5025.1317  21900.86   5305.52362
## Mar 2009 -5391.2925  22007.61   5657.42787
## Apr 2009 -5792.7941  22114.37   5185.92286
## May 2009 -4342.8346  21900.41   3336.17913
## Jun 2009 -747.0787   21686.44    505.63897
## Jul 2009 3294.3252   21406.70 -2959.77100
## Aug 2009 6789.0248   21126.95 -6355.97673
## Sep 2009 8908.2448   21036.99 -7522.73152
## Oct 2009 7890.9851   20947.02 -5446.75664
## Nov 2009 2875.5202   21062.66 -1143.17845
## Dec 2009 -3149.1758  21178.29   2955.88090
## Jan 2010 -5309.7893  21350.19   3734.59544
## Feb 2010 -5025.1317  21522.09   4163.03877
## Mar 2010 -5391.2925  21560.08   4968.71193
## Apr 2010 -5792.7941  21598.07   3643.47582
## May 2010 -4342.8346  21497.63   1451.45609
## Jun 2010 -747.0787   21397.19 -1383.86007
## Jul 2010 3294.3252   21401.78 -4456.10535
## Aug 2010 6789.0248   21406.37 -5952.89639
## Sep 2010 8908.2448   21798.20 -5690.19802
## Oct 2010 7890.9851   22190.03 -3784.76998
## Nov 2010 2875.5202   22960.24 -1039.50804
## Dec 2010 -3149.1758  23730.44   1913.73505
## Jan 2011 -5309.7893  24458.01   2546.78199
## Feb 2011 -5025.1317  25185.57   2774.55772
## Mar 2011 -5391.2925  25467.31   4137.72807
## Apr 2011 -5792.7941  25749.05   5157.48916
## May 2011 -4342.8346  25714.60   5310.73723
## Jun 2011 -747.0787   25680.14   1959.43888
## Jul 2011 3294.3252   25597.47 -2550.54120
## Aug 2011 6789.0248   25514.79 -5410.06704
## Sep 2011 8908.2448   25471.23 -6628.22914
## Oct 2011 7890.9851   25427.68 -5506.16158
## Nov 2011 2875.5202   25518.63 -1464.15347
## Dec 2011 -3149.1758  25609.59   3354.58580
```

##	Jan	2012	-5309.7893	25896.70	4015.58495
##	Feb	2012	-5025.1317	26183.82	2507.56289
##	Mar	2012	-5391.2925	26539.92	1961.37151
##	Apr	2012	-5792.7941	26896.02	1859.27087
##	May	2012	-4342.8346	27214.11	993.72576
##	Jun	2012	-747.0787	27532.19	-1185.11579
##	Jul	2012	3294.3252	27900.64	-3489.96780
##	Aug	2012	6789.0248	28269.09	-5273.11556
##	Sep	2012	8908.2448	28924.57	-4225.31221
##	Oct	2012	7890.9851	29580.04	-438.52919
##	Nov	2012	2875.5202	30539.16	885.31627
##	Dec	2012	-3149.1758	31498.28	1544.64289
##	Jan	2013	-5309.7893	32448.96	1562.07654
##	Feb	2013	-5025.1317	33399.64	100.48899
##	Mar	2013	-5391.2925	34067.22	1169.07426
##	Apr	2013	-5792.7941	34734.79	3625.50026
##	May	2013	-4342.8346	35232.27	3704.31585
##	Jun	2013	-747.0787	35729.74	1639.83502
##	Jul	2013	3294.3252	36289.18	-1763.50248
##	Aug	2013	6789.0248	36848.61	-5000.13572
##	Sep	2013	8908.2448	37505.37	-4353.61903
##	Oct	2013	7890.9851	38162.14	-1846.87266
##	Nov	2013	2875.5202	39090.97	-1238.99406
##	Dec	2013	-3149.1758	40019.81	869.36571
##	Jan	2014	-5309.7893	41425.16	1819.62759
##	Feb	2014	-5025.1317	42830.51	-565.38172
##	Mar	2014	-5391.2925	44613.99	-1821.44659
##	Apr	2014	-5792.7941	46397.46	-832.17072
##	May	2014	-4342.8346	48535.17	-1278.58590
##	Jun	2014	-747.0787	50672.88	-1308.29750
##	Jul	2014	3294.3252	53236.11	-1799.18324
##	Aug	2014	6789.0248	55799.34	-3443.36473
##	Sep	2014	8908.2448	58564.59	-2606.58688
##	Oct	2014	7890.9851	61329.84	-1527.07936
##	Nov	2014	2875.5202	64654.45	-506.22072
##	Dec	2014	-3149.1758	67979.06	2950.11908
##	Jan	2015	-5309.7893	73140.35	1583.19031
##	Feb	2015	-5025.1317	78301.64	-2809.00967
##	Mar	2015	-5391.2925	85495.95	-11415.90859
##	Apr	2015	-5792.7941	92690.26	-19062.46677
##	May	2015	-4342.8346	99738.51	-18679.42489
##	Jun	2015	-747.0787	106786.76	-9163.42943
##	Jul	2015	3294.3252	111991.18	6761.99792
##	Aug	2015	6789.0248	117195.60	22912.87953
##	Sep	2015	8908.2448	120633.84	35617.91227
##	Oct	2015	7890.9851	124072.09	36356.92467
##	Nov	2015	2875.5202	125546.09	21077.13936
##	Dec	2015	-3149.1758	127020.09	-4464.66480
##	Jan	2016	-5309.7893	125961.75	-14198.21533
##	Feb	2016	-5025.1317	124903.42	-10839.53707
##	Mar	2016	-5391.2925	121520.44	-6955.39849
##	Apr	2016	-5792.7941	118137.46	-4953.41918
##	May	2016	-4342.8346	113766.11	859.22025
##	Jun	2016	-747.0787	109394.77	9107.31325



```
## Jul 2016 3294.3252 105054.60 16826.07286
## Aug 2016 6789.0248 100714.44 22281.53671
## Sep 2016 8908.2448 96079.49 11978.52007
## Oct 2016 7890.9851 91444.53 -8428.01689
## Nov 2016 2875.5202 86320.69 -15361.20755
## Dec 2016 -3149.1758 81196.84 -12228.91705
## Jan 2017 -5309.7893 76160.67 -8963.38214
## Feb 2017 -5025.1317 71124.50 -3939.36845
## Mar 2017 -5391.2925 67424.06 -402.77190
## Apr 2017 -5792.7941 63723.63 -63.33462
## May 2017 -4342.8346 61776.12 -677.03061
## Jun 2017 -747.0787 59828.60 -877.77302
## Jul 2017 3294.3252 58709.95 -2058.02039
## Aug 2017 6789.0248 57591.29 -3014.06351
## Sep 2017 8908.2448 56858.01 -4731.25440
## Oct 2017 7890.9851 56124.73 -3610.71563
## Nov 2017 2875.5202 55742.38 -1715.40447
## Dec 2017 -3149.1758 55360.04 580.38785
## Jan 2018 -5309.7893 55053.38 2623.90565
## Feb 2018 -5025.1317 54746.73 1542.15226
## Mar 2018 -5391.2925 54132.26 1294.03268
## Apr 2018 -5792.7941 53517.79 3268.75384
## May 2018 -4342.8346 52991.30 3645.28718
## Jun 2018 -747.0787 52464.80 1993.52409
## Jul 2018 3294.3252 51894.89 -632.96836
## Aug 2018 6789.0248 51324.98 -4564.00655
## Sep 2018 8908.2448 50680.40 -5004.89305
## Oct 2018 7890.9851 50035.81 -213.04987
## Nov 2018 2875.5202 49350.87 3146.10821
```

The additive seasonal coefficients for each month are:

```
Jan -5309.7893
Feb -5025.1317
Mar -5391.2925
Apr -5792.7941
May -4342.8346
Jun -747.0787
Jul 3294.3252
Aug 6789.0248
Sep 8908.2448
Oct 7890.9851
Nov 2875.5202
Dec -3149.1758
```


---

## 4. Is the data stationary?

“If we fit a stationary model to data, we assume our data are a realization of a stationary process. So our first step in an analysis should be to check whether there is any evidence of a trend or seasonal effects and, if there is, remove them.”

```
values.fit.ma <- ts(na.omit(values.fit.ma), frequency=12, start=c(2008, 1))
acf(values.fit.ma)
```

## Series Variance



```
#adf.test(values.fit.ma)
```

The data doesn't appear to be stationary - it is apparent that mean and variance change over time. Looking at the ACF plot we can also reach the same conclusion, given there is no sharp drop in the values, but instead a smooth decay. So we need to transform our data to make it more stationary in order to use it for ARIMA models, for instance. This is further confirmed by running the ADF test: with a p-value = 0.7279 we can't reject the null hypothesis of non-stationarity.

```
adf.test(values.fit.ma)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: values.fit.ma
## Dickey-Fuller = -1.6359, Lag order = 5, p-value = 0.7279
## alternative hypothesis: stationary
```


### Making the data stationary

Some of the possible mathematical transforms include: differencing, log (and Box-Cox), moving average, percent change, lag, or cumulative sum.

#### Differencing

seasonal differencing


```
values_diff12 = diff(values.fit.ma, lag = 12)
tm <- cbind(values, values_diff12)
plot(tm)
```



trend

remains, take 1st difference


```
values_diff12_1 = diff(values_diff12)
tm <- cbind(values, values_diff12, values_diff12_1)
plot(tm)
```



A first difference seems to suffice here, but let's check again to confirm the visual inspection by using an ACF plot

and ADF test:

```
acf(na.omit(values_diff12_1), main="ACF for series after 1st difference")
```





```
adf.test(na.omit(values_diff12_1))
```

```
##
## Augmented Dickey-Fuller Test
##
## data: na.omit(values_diff12_1)
## Dickey-Fuller = -3.5973, Lag order = 4, p-value = 0.03656
## alternative hypothesis: stationary
```

Both the ACF plot and ADF test confirm non-stationarity after the first difference. There's a sort of pattern going on on this ACF plot, probably due to the seasonal component - not a pure SMA process, indicates AR.

```
diff.values <- na.omit(values_diff12_1)
par(mfrow=c(2,1))
acf(diff.values, main="")
pacf(diff.values, main="")
```



A relatively slow decay until lag 6, might indicate a AR of order 6. Negative spike at lag 1 is possible SAR(1) term, it repeats.




## 5. Model

### SARIMA(p, d, q)(P, D, Q) order

Try for up to 1 SAR, 1 difference, 1 seasonal D, 1-6 AR, review ACF and PACF if no good fits found.

```
#start with SAR(1)
values.fit.sarima1 <- arima(values.fit.ma, order=c(1, 1, 0), seasonal = list(order = c(1, 1, 0), period
values.fit.sarima1

##
## Call:
## arima(x = values.fit.ma, order = c(1, 1, 0), seasonal = list(order = c(1, 1,
##      0), period = 12))
##
## Coefficients:
##          ar1      sar1
##      0.7262  -0.4014
## s.e.  0.0630   0.0797
##
## sigma^2 estimated as 25281493:  log likelihood = -1164.6,  aic = 2335.21
#residuals
tsdisplay(residuals(values.fit.sarima1), lag.max=45, main='SARIMA Model 1 Residuals')
```



```
#forecast accuracy
```

```
forecast.sarima1 <- forecast(values.fit.sarima1, h=12) #12months
accuracy(forecast.sarima1)
```

```
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 4.127787 4770.048 2765.512 0.2906328 4.655959 0.1456898
##               ACF1
## Training set 0.4474625
```




```
#1 ar wasnt good, look at 3
```

```
values.fit.sarima2 <- arima(values.fit.ma, order=c(3, 1, 0), seasonal = list(order = c(1, 1, 0), period = 12))
values.fit.sarima2
```

```
##
## Call:
## arima(x = values.fit.ma, order = c(3, 1, 0), seasonal = list(order = c(1, 1, 0), period = 12))
##
## Coefficients:
##      ar1      ar2      ar3      sar1
##      1.3649 -0.9807 0.2998 -0.4206
## s.e.  0.0895  0.1283 0.0898  0.0816
##
## sigma^2 estimated as 14146123:  log likelihood = -1131.38,  aic = 2272.76
```

```
#residuals
```

```
tsdisplay(residuals(values.fit.sarima2), lag.max=45, main='SARIMA Model 2 Residuals')
```



Lag

Lag

*#forecast accuracy*

```
forecast.sarima2 <- forecast(values.fit.sarima2, h=12) #12months
accuracy(forecast.sarima2)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -58.74926 3568.127 2171.997 0.1388934 3.892016 0.1144229
##              ACF1
## Training set 0.1122667
```

Results aren't very good




*#ar6*

```
values.fit.sarima3 <- arima(values.fit.ma, order=c(6, 1, 0), seasonal = list(order = c(1, 1, 0), period
values.fit.sarima3
```

```
##
## Call:
## arima(x = values.fit.ma, order = c(6, 1, 0), seasonal = list(order = c(1, 1,
##    0), period = 12))
##
## Coefficients:
##      ar1      ar2      ar3      ar4      ar5      ar6      sar1
##    1.6753 -1.8012 1.5463 -1.2726 0.7647 -0.2654 -0.4099
## s.e. 0.0894 0.1658 0.2026 0.2037 0.1685 0.0915 0.0818
##
## sigma^2 estimated as 10144043: log likelihood = -1112.64, aic = 2241.27
```

*#residuals*

```
tsdisplay(residuals(values.fit.sarima3), lag.max=45, main='SARIMA Model 3 Residuals')
```



```
#forecast accuracy
forecast.sarima3 <- forecast(values.fit.sarima3, h=12) #12months
accuracy(forecast.sarima3)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -61.5572 3021.532 1750.562 0.1144758 3.102899 0.09222128
##              ACF1
## Training set 0.1011675
```

## Auto.arima


Another method is to let auto.arima estimate parameters and see how they fit with our conclusions so far:

```
values.fit.sarima4 <- auto.arima(ma.fit, seasonal = TRUE)
values.fit.sarima4
```

```
## Series: ma.fit
## ARIMA(5,1,0)(1,0,0)[12]
##
## Coefficients:
##      ar1      ar2      ar3      ar4      ar5      sar1
##      1.6199 -1.6228  1.2198 -0.7996  0.2690  0.1627
## s.e.  0.0855  0.1512  0.1786  0.1520  0.0862  0.0898
##
## sigma^2 estimated as 7622202:  log likelihood=-1203.79
## AIC=2421.57  AICc=2422.5  BIC=2441.59
```

```
#residuals
tsdisplay(residuals(values.fit.sarima4), lag.max=45, main='SARIMA Model 4 Residuals')
```





So, our initial interpretation of PACF was wrong...auto suggests 5 non-seasonal terms and 1 seasonal for AR. . Let's keep this auto model and check for other types of models


## Holt Winters

Trying a different model, exponential smoothing

```
values.fit.hw1 <- ets(values.fit.ma, model="MAM", damped=FALSE)
values.fit.hw1
```

```
## ETS(M,A,M)
##
## Call:
## ets(y = values.fit.ma, model = "MAM", damped = FALSE)
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.1047
##   gamma = 1e-04
##
## Initial states:
##   l = 18282.9878
##   b = 176.2215
##   s = 0.9391 0.9758 1.0617 1.1418 1.1265 1.0693
##       1.0262 0.9741 0.9228 0.9095 0.921 0.9322
##
## sigma: 0.0568
##
##      AIC      AICc      BIC
## 2656.877 2662.341 2705.625
```

```
plot(values.fit.hw1)
```



```
accuracy(values.fit.hw1)
```

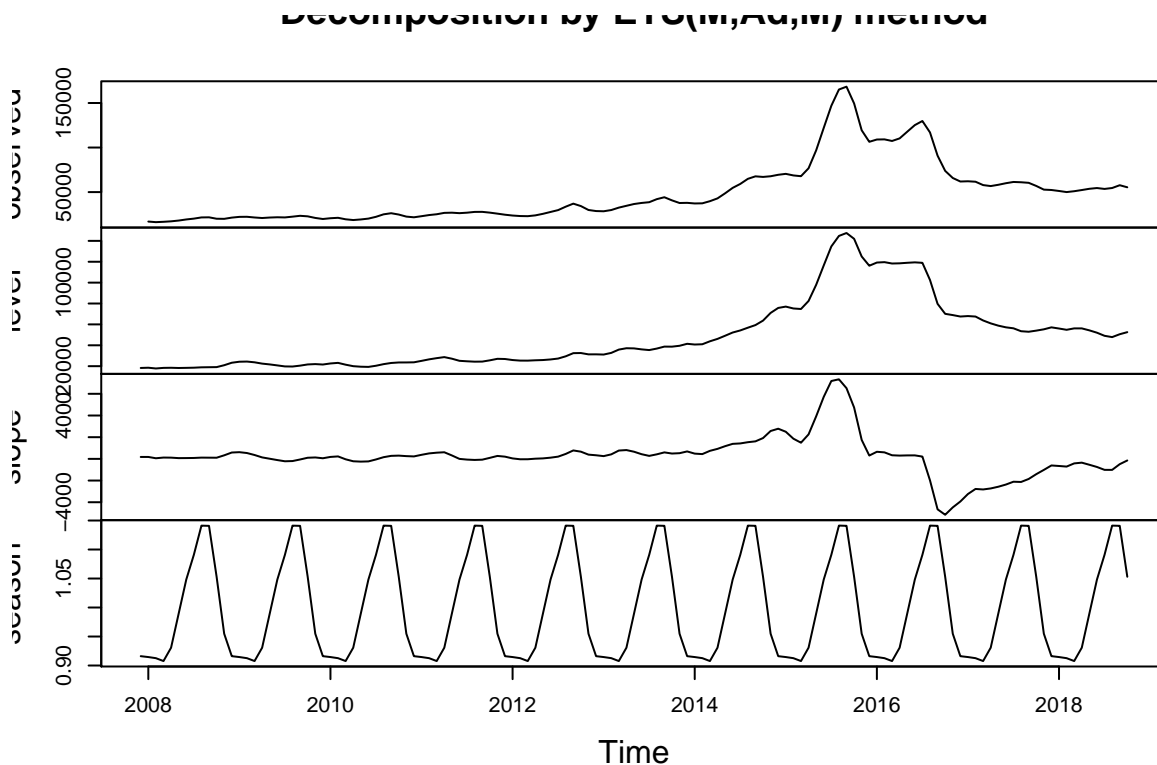
```
##           ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -105.8016 4836.084 2297.052 0.2373048 3.912748 0.1210109
##           ACF1
## Training set 0.6923798
```

```
values.fit.hw2 <- ets(values.fit.ma, model="MAM",
                      damped=TRUE)
values.fit.hw2
```

```
## ETS(M,Ad,M)
##
## Call:
## ets(y = values.fit.ma, model = "MAM", damped = TRUE)
##
## Smoothing parameters:
##   alpha = 0.9998
##   beta  = 0.1311
##   gamma = 2e-04
##   phi   = 0.9647
##
## Initial states:
##   l = 18282.7484
##   b = 176.2974
##   s = 0.9161 0.9548 1.053 1.1408 1.1412 1.0908
##       1.0482 0.9895 0.9307 0.9077 0.9127 0.9146
##
```

```
## sigma: 0.0561
##
##      AIC      AICc      BIC
## 2654.275 2660.437 2705.890
```

```
plot(values.fit.hw2)
```



```
accuracy(values.fit.hw2)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 26.0826 4468.495 2220.006 0.3683582 3.867666 0.116952
##              ACF1
## Training set 0.6704183
```

pick the model with damped trend, hw2, marginally better

## 6. Model evaluation and forecasting

### Accuracy

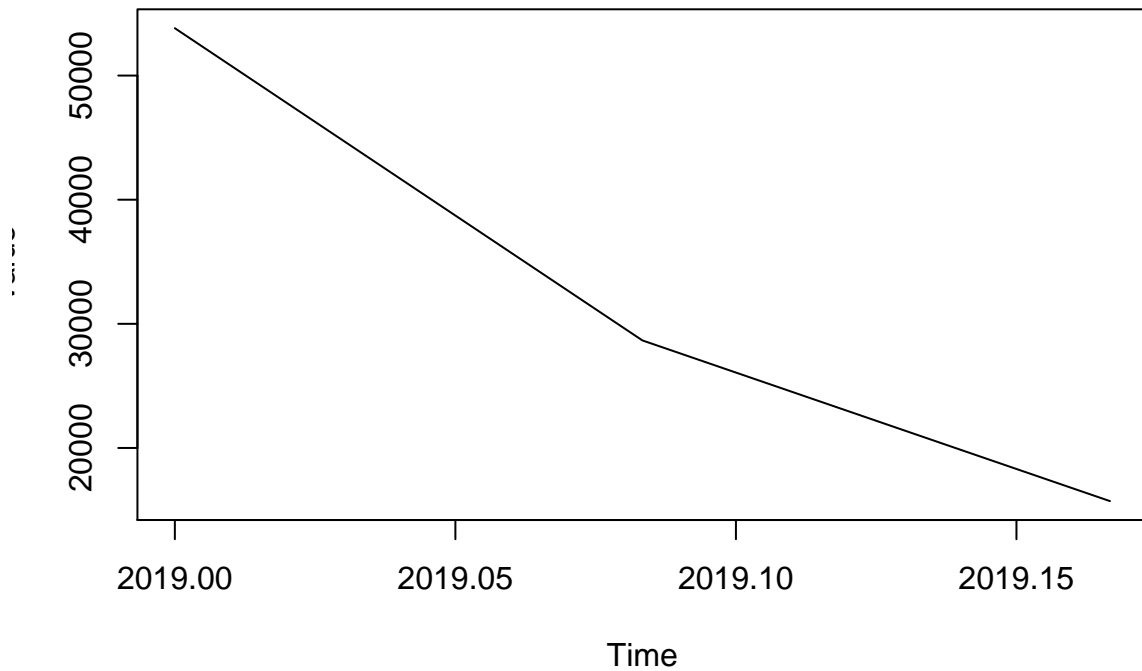
Compare with real observed values for the first three months of 2019

```
data2019 <- read_csv("migr_asylum_applicant_EU28/migr_asyappctzm_2019_Data.csv")
```

```
## Parsed with column specification:
## cols(
##   TIME = col_character(),
##   GEO = col_character(),
##   CITIZEN = col_character(),
##   SEX = col_character(),
```

```
## AGE = col_character(),
## ASYL_APP = col_character(),
## UNIT = col_character(),
## Value = col_number(),
## `Flag and Footnotes` = col_character()
## )

values2019 <- ts(data2019[8], start = c(2019, 1), frequency=12)
plot.ts(values2019)
```



downward trend

A strong


## Forecasting using the chosen model(s)

Forecasting the next 12 months

## Seasonal ARIMA

```
forecast.values.sarima4 <- forecast(values.fit.sarima4, h=12)
plot(forecast.values.sarima4)
lines(values2019, col="cyan", lwd=2, lty=3)
```

Forecasts from ARIMA(0,1,0)(1,0,0)[12]



```
accuracy(forecast.values.sarima4, values2019)
```

```
##              ME      RMSE      MAE      MPE      MAPE
## Training set   56.51262 2685.475 1659.178  0.4042005  3.198991
## Test set      -11288.61921 19383.377 17658.847 -71.8985574 83.735829
##              MASE      ACF1 Theil's U
## Training set  0.08740708  0.0901048      NA
## Test set      0.93028488 -0.0173775  1.589794
```

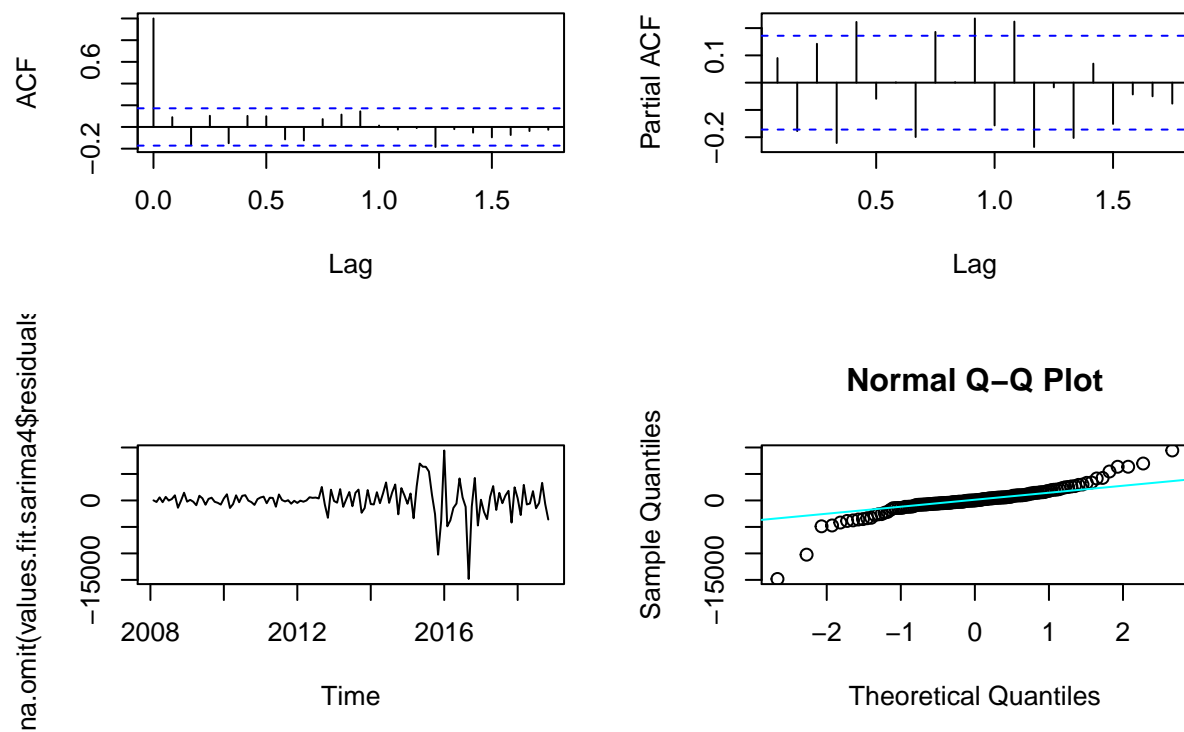
```
#test residuals
```

```
Box.test(residuals(values.fit.sarima4), type="Ljung-Box") #test alternative hypothesis that there is no
```

```
##
## Box-Ljung test
##
## data: residuals(values.fit.sarima4)
## X-squared = 1.08, df = 1, p-value = 0.2987
```

```
par(mfrow=c(2,2))
acf(na.omit(values.fit.sarima4$residuals))
pacf(na.omit(values.fit.sarima4$residuals))
plot(na.omit(values.fit.sarima4$residuals))
qqnorm(na.omit(values.fit.sarima4$residuals))
qqline(na.omit(values.fit.sarima4$residuals), col="cyan")
```


series na.omit(values.fit.sarima4\$residuals) series na.omit(values.fit.sarima4\$residuals)



## Holt Winters seasonal smoothing

```
plot(forecast(values.fit.hw2))
lines(values2019, col="cyan", lwd=2, lty=3)
```

Forecasts from ETS(m,Ad,m)



```
accuracy(forecast(values.fit.hw2), values2019)
```

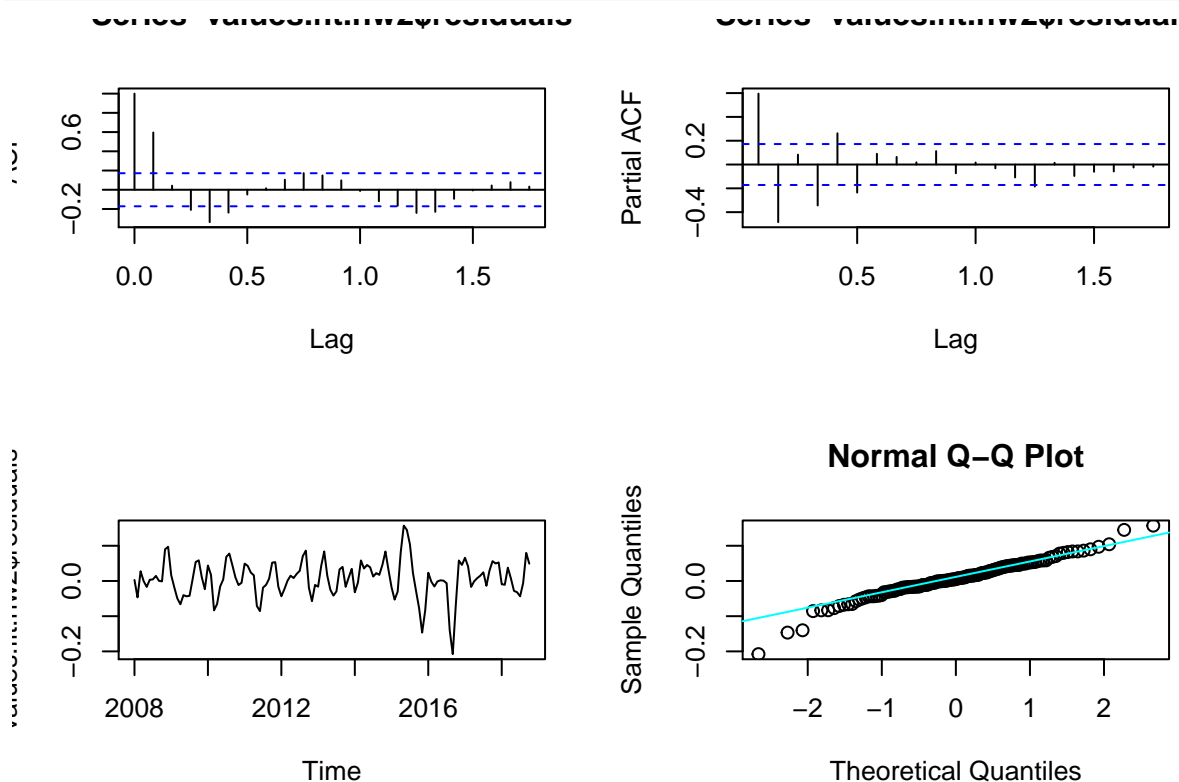
```
##           ME      RMSE      MAE      MPE      MAPE
## Training set  26.0826 4468.495 2220.006  0.3683582  3.867666
## Test set     -14691.8277 21416.809 18769.668 -84.6273976 92.204913
##           MASE      ACF1 Theil's U
## Training set  0.116952  0.67041833      NA
## Test set     0.988804 -0.02329674  1.768313
```

```
#test residuals
```

```
Box.test(residuals(values.fit.hw2), type="Ljung-Box") #test alternative hypothesis that there is non-zero autocorrelation
```

```
##
## Box-Ljung test
##
## data: residuals(values.fit.hw2)
## X-squared = 47.177, df = 1, p-value = 6.485e-12
```

```
par(mfrow=c(2,2))
acf(values.fit.hw2$residuals)
pacf(values.fit.hw2$residuals)
plot(values.fit.hw2$residuals)
qqnorm(values.fit.hw2$residuals)
qqline(values.fit.hw2$residuals, col="cyan")
```



great on the residuals correlations accuracy is bad on this test

```
forecast.values.sarima4
```

```
##           Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Dec 2018      45451.82 35529.760 55373.87 30277.3405 60626.29
## Jan 2019      44259.66 28054.117 60465.20 19475.4218 69043.89
```

not

## Feb 2019	43551.34	22416.156	64686.52	11227.8662	75874.81
## Mar 2019	44254.86	19334.940	69174.78	6143.1303	82366.59
## Apr 2019	46249.75	18557.065	73942.43	3897.4433	88602.06
## May 2019	47719.25	18110.435	77328.07	2436.4733	93002.04
## Jun 2019	47771.61	16640.342	78902.87	160.4459	95382.77
## Jul 2019	47060.43	14381.062	79739.80	-2918.3507	97039.21
## Aug 2019	46712.42	12322.702	81102.14	-5882.1142	99306.95
## Sep 2019	46659.28	10441.561	82877.00	-8730.9412	102049.50
## Oct 2019	45938.72	7842.294	84035.15	-12324.7356	104202.18
## Nov 2019	45058.51	4956.470	85160.55	-16272.2678	106389.29